

Liquid PureScript Feasibility Study

Status: active **Category:** research **Created:** 2026-01-29 **Author:** Claude (research session)

Executive Summary

Building a Liquid Haskell-style refinement type system for PureScript is **technically feasible** but represents a significant engineering effort. The key enablers are:

1. **CoreFn intermediate representation** - PureScript already outputs a well-documented IR
2. **purescript-z3 bindings** - Z3 SMT solver is already accessible from PureScript
3. **Backend-agnostic architecture** - Multiple backends (JS, Erlang, etc.) consume CoreFn

The main challenges are:

1. No compiler plugin system (would need to fork compiler or build external tool)
2. Type erasure complications (which refinements survive to runtime?)
3. FFI boundary handling (foreign code can't be verified)

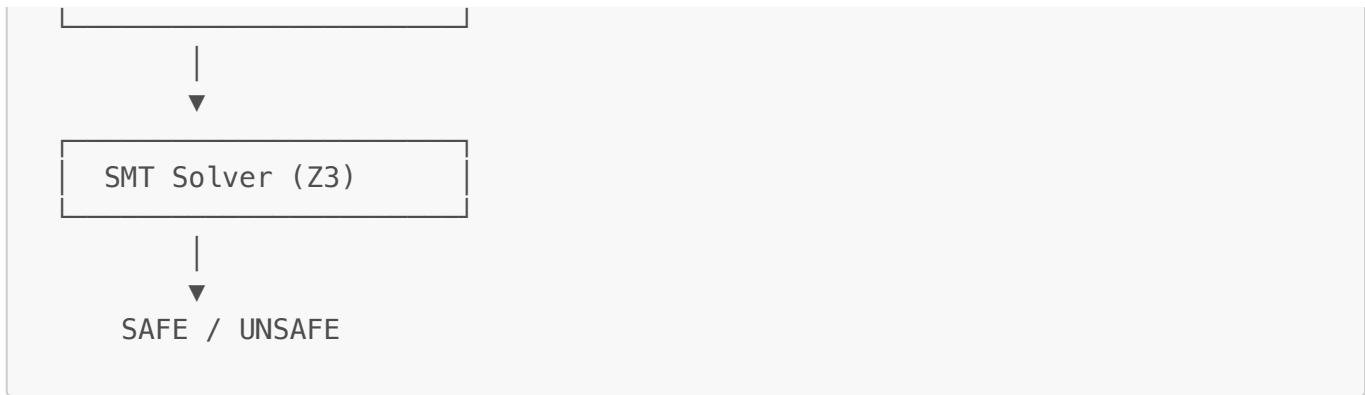
Recommended approach: Build as an external tool consuming CoreFn, not as a compiler modification.

Part 1: How Liquid Haskell Works

Architecture Overview

Liquid Haskell operates as a GHC plugin, hooking into the `typecheckResultAction` phase:





Key Technical Choices

1. **Unoptimized Core:** LH requires `-O0` because optimizations can change type representations
2. **Horn Clauses:** Verification conditions are encoded as Horn clauses for SMT
3. **Annotation Syntax:** Refinements live in specially-formatted Haskell comments
4. **Decidable Logic:** Uses QF-EUFLIA (quantifier-free equality, uninterpreted functions, linear integer arithmetic)

What LH Verifies

```

{-@ type Pos = { v:Int | v > 0 } @-}
{-@ abs :: Int -> Pos @-}
abs :: Int -> Int
abs n | n > 0      = n
      | otherwise = -n
  
```

LH generates constraints:

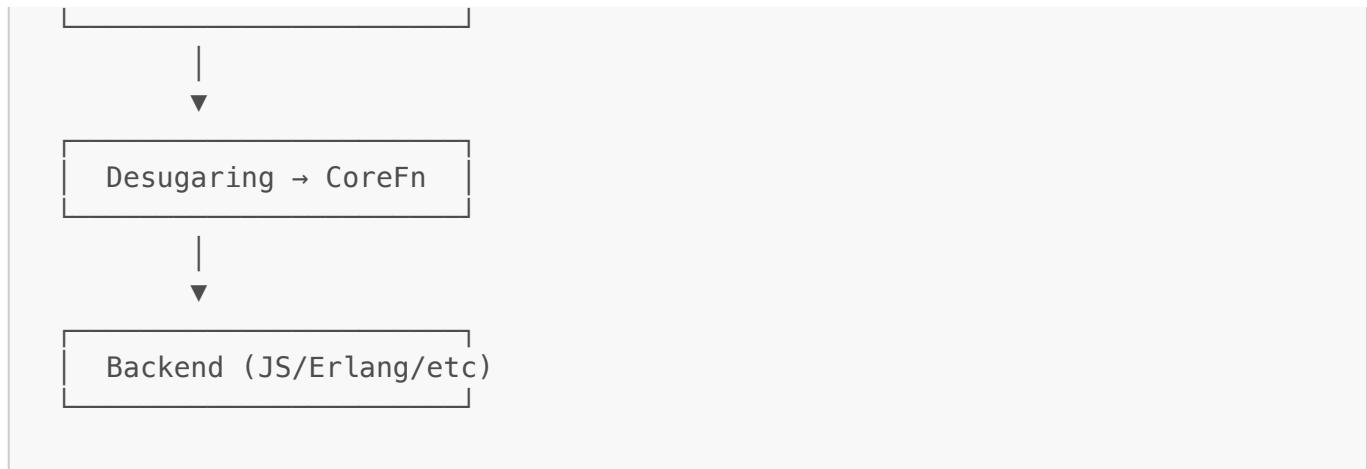
- If `n > 0`, prove `n > 0` (trivial)
- If `not (n > 0)`, prove `-n > 0` (follows from `n <= 0 ∧ n ≠ 0` for non-zero inputs)

Part 2: PureScript's Compiler Architecture

CoreFn: The Key Enabler

PureScript compiles to CoreFn, a simplified functional IR:





CoreFn is:

- Well-documented and stable
- Dumped as JSON (`purs compile --codegen corefn`)
- Used by multiple backends (purerl, purescript-backend-optimizer)
- Simpler than source (no type classes, no do-notation sugar)

No Plugin System

Unlike GHC, PureScript doesn't have a compiler plugin system. Options:

1. **Fork the compiler** - Invasive, maintenance burden
2. **External tool consuming CoreFn** - Clean, independent development
3. **Build into Spago** - Would still need to consume CoreFn

Recommendation: External tool is the pragmatic path.

Type Information in Externs

`externs.json` files contain exported type signatures. Combined with CoreFn, this provides enough information for refinement checking.

Part 3: Existing PureScript Building Blocks

purescript-z3

[purescript-z3](#) provides FFI bindings to Z3:

```

import Z3 as Z3

solve :: Effect (Maybe Solution)
solve = Z3.run do
  x <- Z3.int "x"
  y <- Z3.int "y"
  Z3.assert (x `Z3.gt` Z3.intValue 0)
  Z3.assert (y `Z3.eq` (x `Z3.mul` Z3.intValue 2))
  Z3.withModel \model -> do
    xVal <- Z3.eval model x
  
```

```
yVal <- Z3.eval model y
pure { x: xVal, y: yVal }
```

Status: Partial bindings, but sufficient for constraint solving. Uses Z3's WASM build.

purescript-refined

[purescript-refined](#) provides runtime-checked refinement types:

```
type DiceRoll = Refined (FromTo D1 D6) Int

validRoll :: Either RefinedError DiceRoll
validRoll = refine 5 -- Right DiceRoll

invalidRoll :: Either RefinedError DiceRoll
invalidRoll = refine 8 -- Left (FromToError 1 6 8)
```

Limitation: Runtime checking only. A Liquid system would verify these statically.

purescript-corefn

[purescript-corefn](#) provides PureScript types for working with CoreFn JSON:

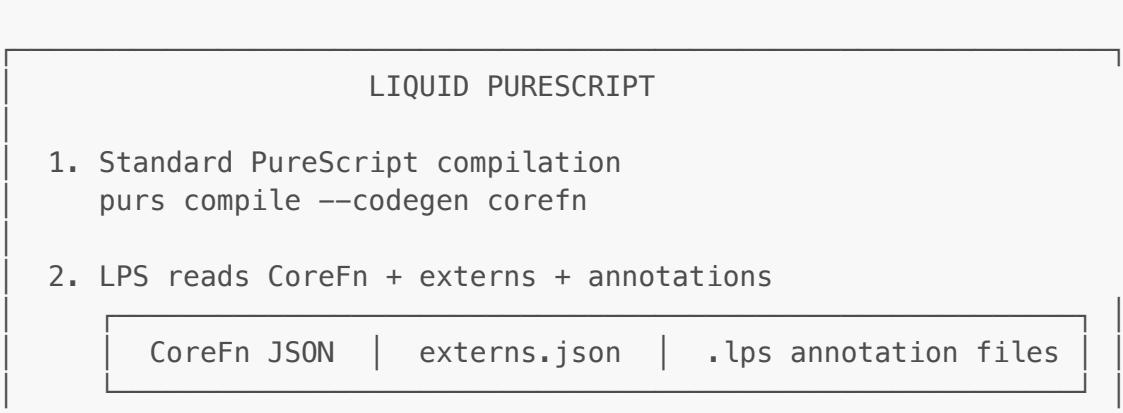
```
import PureScript.CoreFn (Module, readModuleJSON)

loadModule :: String -> Effect Module
loadModule path = do
  json <- readTextFile path
  pure (readModuleJSON json)
```

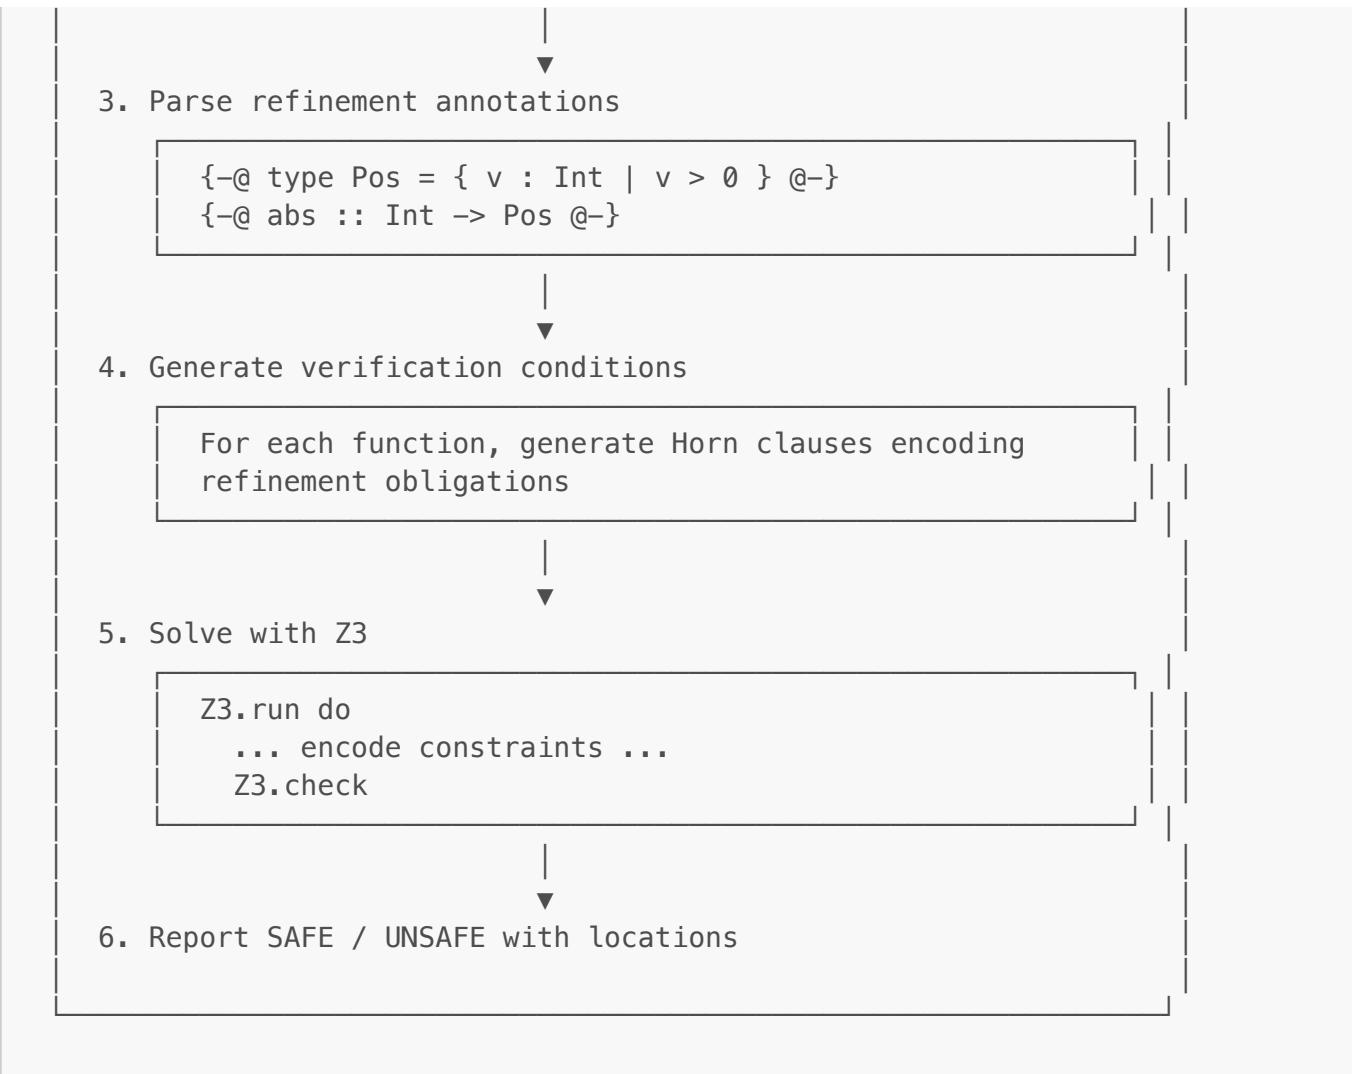
Status: Useful for parsing CoreFn, though may need updates for recent compiler versions.

Part 4: Architecture for Liquid PureScript

Proposed Pipeline

- 
1. Standard PureScript compilation
purs compile --codegen corefn
 2. LPS reads CoreFn + externs + annotations

CoreFn JSON		externs.json		.lps annotation files
-------------	--	--------------	--	-----------------------



Annotation Format

Following LH's lead, annotations in specially-marked comments:

```

-- .purs file
module MyModule where

{-@ type NonEmpty a = { v : Array a | length v > 0 } @-}

{-@ head :: NonEmpty a -> a @-}
head :: forall a. Array a -> a
head xs = unsafePartial (Array.head xs)

{-@ filter :: (a -> Boolean) -> Array a -> Array a @-}
filter :: forall a. (a -> Boolean) -> Array a -> Array a
filter = Array.filter

-- This would be UNSAFE - filter might return empty array
{-@ badHead :: Array Int -> Int @-}
badHead xs = head (filter (_ > 0) xs) -- ERROR: Cannot prove non-empty
  
```

Alternative: Separate Spec Files

Like LH's `.spec` files:

```
-- MyModule.lps
module MyModule where

type NonEmpty a = { v : Array a | length v > 0 }
type Pos = { v : Int | v > 0 }

head :: NonEmpty a -> a
abs :: Int -> Pos
```

Advantage: No changes to `.purs` files, cleaner separation.

Part 5: Technical Challenges

Challenge 1: Type Erasure

PureScript erases types at runtime. For refinements:

```
{-@ type Pos = { v : Int | v > 0 } @-}

{-@ add :: Pos -> Pos -> Pos @-}
add :: Int -> Int -> Int
add x y = x + y
```

At runtime, `Pos` is just `Int`. The refinement exists only at verification time. This is fine for pure code but problematic for FFI boundaries.

Solution: Treat FFI as trust boundary. Foreign functions get assumed types.

Challenge 2: Higher-Order Functions

Refinements on function arguments:

```
{-@ map :: (a -> b) -> Array a -> Array b @-}
{-@ map :: (Pos -> Pos) -> NonEmpty Pos -> NonEmpty Pos @-} -- More
specific
```

LH handles this with "abstract refinements" - refinement variables that can be instantiated.

Complexity: Significant. This is where LH research papers get deep.

Challenge 3: Records and Row Types

PureScript's row-polymorphic records are more flexible than Haskell's:

```
type Person r = { name :: String, age :: Int | r }
```

How do refinements interact with row polymorphism?

```
{-@ type Adult r = { name :: String, age :: { v : Int | v >= 18 } | r }
@-}
```

Status: Novel research territory. No existing system handles this well.

Challenge 4: Type Classes

Refinements that vary by instance:

```
class Monoid a where
  mempty :: a
  mappend :: a -> a -> a

-- For Array: mempty = [] (length 0)
-- For NonEmptyArray: mempty = [default] (length 1)
```

LH Approach: Bounded refinement types, instance-specific specs.

Challenge 5: Laziness

Haskell is lazy; LH has complex handling for divergence. PureScript is strict, which simplifies things:

```
-- In strict PureScript, this always diverges:
loop :: forall a. a
loop = loop

-- In lazy Haskell, it only diverges if forced
```

Advantage: PureScript's strictness makes verification simpler.

Part 6: Implementation Roadmap

Phase 1: Minimal Viable Verifier (4-6 weeks)

Goal: Verify refinements on simple functions (no HOF, no polymorphism)

Deliverables:

1. CoreFn parser (or update purescript-corefn)
2. Annotation parser (comment format or separate files)

3. VC generator for basic expressions
4. Z3 integration for constraint solving
5. Error reporting with source locations

Demo: Verify `abs :: Int -> Pos` style functions

Phase 2: Polymorphism and Records (6-8 weeks)

Goal: Handle parametric polymorphism and records

Deliverables:

1. Abstract refinements for polymorphism
2. Record refinement syntax
3. Measure definitions (e.g., `length` for arrays)
4. Pre/post conditions on record fields

Demo: Verify `head :: NonEmpty a -> a`

Phase 3: Higher-Order Functions (8-12 weeks)

Goal: Verify functions taking function arguments

Deliverables:

1. Refinement inference for lambdas
2. Abstract refinement instantiation
3. Contravariant argument handling
4. Fix point inference for recursive functions

Demo: Verify `map`, `filter`, `fold` with refinements

Phase 4: Type Classes (Research)

Goal: Instance-specific refinements

Deliverables:

1. Bounded refinement types
2. Instance specs
3. Coherent refinement inheritance

Status: This requires genuine research, not just engineering.

Part 7: Comparison with Alternatives

Option A: Liquid PureScript (this document)

Approach: External tool, SMT-based verification **Pros:** Proven approach (LH), automatic verification, decidable **Cons:** Complex implementation, limited expressiveness

Option B: Full Dependent Types in Compiler

Approach: Modify PureScript compiler to support dependent types **Pros:** Full expressiveness, no separate tool **Cons:** Massive compiler change, breaks backward compatibility

From [GitHub issue #2214](#):

"implement a typechecker that supports dependent types along with records and typeclasses" – described as "the trivial task" (sarcastically)

Status: Closed as out-of-scope research project.

Option C: Idris Backend for JavaScript

Approach: Use Idris (dependently typed language) with JS backend **Pros:** Already exists, full dependent types **Cons:** Different language, immature JS backend

Option D: purescript-refined (Runtime Checking)

Approach: Library-based runtime checking **Pros:** Already works, no compiler changes **Cons:** Runtime cost, no static guarantees

Recommendation

Start with Option A (Liquid PureScript) because:

1. Proven approach from Liquid Haskell
2. External tool = independent development
3. Automatic verification = lower barrier for users
4. Can target important subset (non-empty arrays, positive numbers, etc.)

Part 8: Integration with Playground Vision

The typed feedback loop from the playground research becomes even more powerful with refinements:

CODE	REFINEMENTS
<code>xs = [1, 2, 3]</code>	<code>xs : NonEmpty Int</code> ✓ <code>length = 3 > 0</code>
<code>ys = filter (_ > 5) xs</code>	<code>ys : Array Int</code> △ <code>might be empty</code>
<code>z = head ys</code>	<code>z : ERROR</code> ✗ <code>head requires NonEmpty</code>
<code>-- Fix: add guard</code> <code>z = case ys of</code> <code>[] -> Nothing</code> <code>_ -> Just (head ys)</code>	<code>z : Maybe Int</code> ✓ <code>pattern handles empty</code>

The playground doesn't just show types - it shows **what properties the types guarantee**.

Resources

Liquid Haskell

- [LiquidHaskell Homepage](#)
- [LH Tutorial](#)
- [LH as GHC Plugin \(Well-Typed\)](#)
- [LH ICFP'14 Paper \(PDF\)](#)

PureScript Internals

- [CoreFn source](#)
- [purescript-corefn library](#)
- [purescript-backend-optimizer](#)
- [Dependent Types Discussion](#)

SMT and Z3

- [purescript-z3](#)
- [Z3 JavaScript API](#)
- [SMT-LIB Standard](#)

Related Work

- [Refinement Types for TypeScript \(PDF\)](#)
 - [purescript-refined](#)
-

Conclusion

Liquid PureScript is feasible. The hardest part isn't "can we do it" but "where do we stop":

Scope	Effort	Value
Simple numeric refinements	4-6 weeks	Catch div-by-zero, bounds errors
Non-empty arrays	6-8 weeks	Eliminate partial function errors
Full polymorphism	8-12 weeks	General safety guarantees
Type classes	Research	Instance-specific invariants

The sweet spot is probably **Phase 2** - enough to catch real bugs (non-empty arrays, positive numbers, valid indices) without the research-level complexity of full dependent types.

Combined with the playground vision, this would give PureScript developers something no other JavaScript-targeting language has: **a tight feedback loop with static guarantees about data properties**.